

# NLP: A Super-Simple Introduction


December 5, 2022

# Goals for today


1. Understand the two primary models of language, their assumptions, and when to use them
2. Conceptual understanding of how word embeddings work
3. Conceptual understanding of how embedding models are applied
4. Practical guidance for application and demonstration

Objective

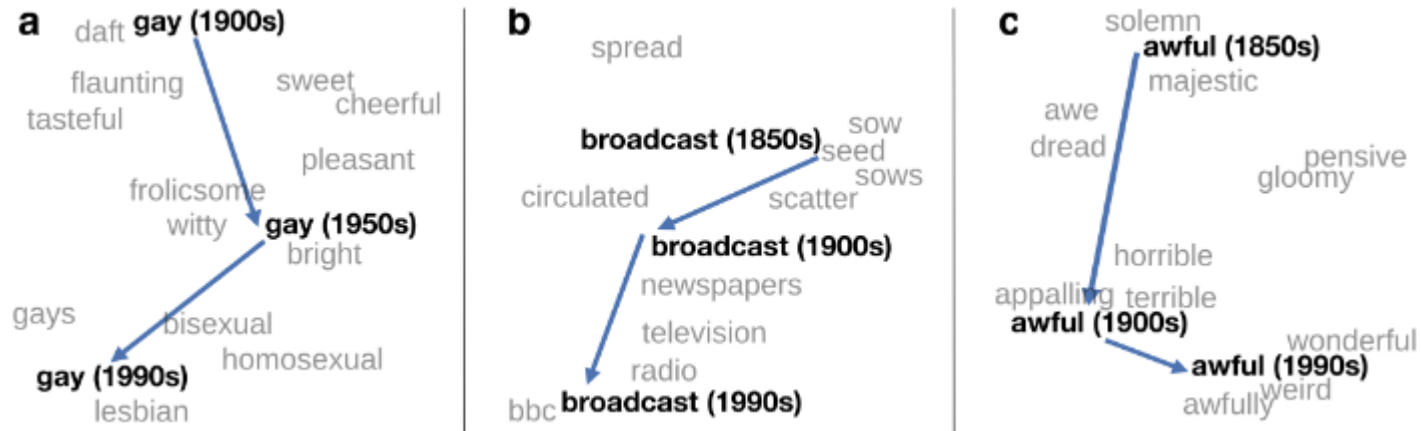
# Objective: Word Math



# Objective: Word Math




# Objective: Word Math



# Approaches

# “Text as Data”




# “NLP”





# Bag-of-Words

- AKA Word counts
- Assumes different types of documents, people, etc. will have different distributions of word use frequency.
- Mostly document or corpus level



# Bag-of-Words: Advantages

- Very fast
- Descriptive analysis
- Very large documents
- Teaches manual text processing and document cleaning

# Bag-of-Words: Disadvantages

- Mostly ignores context
- Generally poor model of semantics
- Generally bad for short documents

$$y_{ijt} \sim \text{Poisson}(\lambda_{ijt})$$
$$\lambda_{ijt} = \exp(\alpha_{it} + \psi_j + \beta_j * \omega_{it})$$



# Bag-of-Words: Disadvantages



# Embeddings


- AKA Distributional semantics, distributional methods, word embeddings, word vectors
- Each word (token) is represented by an  $n$  dimensional vector.
- The value for each dimension is determined by which words are commonly found next to the word in a  $k$  sized window.
- Word level (generally)

# Embeddings

	Animal	Fluffy	Pet	Dangerous	Food
Cat	.98	.94	.87	.30	.15
Dog	.98	.87	.95	.35	.15
Pig	.95	.20	.45	.26	.60
Carrot	.01	.02	.01	.01	.97
Rock	.01	.02	.43	.25	.01

# Embeddings: Advantages

- More theoretically robust model of language
- Far better at semantic representation
- Excellent with small texts
- Almost entirely eliminate pre-processing



# Embeddings: Disadvantages

- Much more computationally intensive
- Document size limits
- Black box
- Neural networks can be hard




# Practical advice



- Use embedding methods as a default
- Bag of Words is still very useful for descriptive analysis
- Bag of Words is still useful for particularly long documents
- Refer to the computer science literature
- Co-author!

# Embeddings

# Word2Vec

“The cat sat on my lap”





```
[48]: ptrump_model.wv['cat']  
[48]: array([ 0.20400783, -0.039106 ,  0.01856888,  0.12811816, -0.19469248,  
              0.0769702 ,  0.20283297,  0.29008055,  0.02028482, -0.12926641,  
             -0.04846543, -0.05207671, -0.03316664, -0.1765712 ,  0.01017776,  
              0.2708853 , -0.08717136,  0.02070806, -0.01205015,  0.14926407,  
              0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568 ,  
             -0.0449645 , -0.09630641, -0.19584368,  0.06084952,  0.02816491,  
             -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,  
             -0.03490731,  0.00283731,  0.15054679,  0.4269414 ,  0.10274042,  
             -0.04721209,  0.1454485 ,  0.24609394, -0.15863296, -0.09804205,  
             -0.17660844,  0.1483835 , -0.15724996, -0.01741939,  0.06103354,  
             -0.08940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,  
             -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,  
              0.26531997,  0.07995594,  0.04015322,  0.3167902 , -0.05698166,  
             -0.2766356 ,  0.14080222,  0.07168636, -0.26628318, -0.13518295,  
              0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,  
              0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,  
              0.0483027 , -0.11918885,  0.04418642, -0.09205839,  0.03029672,  
              0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,  
             -0.08515656, -0.19164322,  0.006798 ,  0.04576558,  0.16265512,  
             -0.19076294,  0.37361056, -0.01433535, -0.09413207, -0.12112152],  
              dtype=float32)
```

```
In [25]: model.similar_by_vector(model['Obama'] - model['USA'] + model['France'])
```


```
Out[25]: [('Sarkozy', 0.704483687877655),  
          ('Obama', 0.7015002369880676),  
          ('President_Nicolas_Sarkozy', 0.642586350440979),
```

```
In [15]: model.similar_by_vector(model['Dublin'] - model['Ireland'] + model['France'])
```

```
Out[15]: [('Paris', 0.740702748298645),  
          ('France', 0.6797398924827576),  
          ('Issy_les_Moulineaux', 0.6246635317802429),
```

# Cosine similarity

- Computing the similarity between words and documents is foundational to everything we do in text analysis
  - Classification
  - Clustering
  - Description
- Cosine (orientation)
- Euclidean (magnitude)
- We use cosine because it is more robust to differences in document size



# The problem with cats

- “This is my cat Popkins”
- “Who let the cat out of the bag?”
- “Joe is one cool cat”
- “Cats are powerful machines”
- “It’s raining cats and dogs”

# Global embeddings



=



```
[48]: ptrump_model.wv['cat']
```

```
[48]: array([ 0.20400783, -0.039106,  0.01856888,  0.12811816, -0.19469248,
            0.0769702,  0.20283297,  0.29008055,  0.02028482, -0.12926641,
           -0.04846543, -0.05207671, -0.03316664, -0.1765712,  0.01017776,
            0.2708853, -0.08717136,  0.02070806, -0.01205015,  0.14926407,
            0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568,
           -0.0449645, -0.09630641, -0.19584368,  0.06084952,  0.02816491,
           -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,
           -0.03490731,  0.00283731,  0.15054679,  0.4269414,  0.10274042,
           -0.04721209,  0.1454485,  0.24609394, -0.15863296, -0.09804205,
           -0.17660844,  0.1483835, -0.15724996, -0.01741939,  0.06103354,
           -0.08940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,
           -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,
            0.26531997,  0.07995594,  0.04015322,  0.3167902, -0.05698166,
            0.2766356,  0.14080222,  0.07168636, -0.26628318, -0.13518295,
            0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,
            0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,
            0.0483027, -0.11918885,  0.04418642, -0.09205839,  0.03029672,
            0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,
           -0.08515656, -0.19164322,  0.006798,  0.04576558,  0.16265512,
           -0.19076294,  0.37361056, -0.01433535, -0.09413207, -0.12112152],
          dtype=float32)
```



# Contextual embeddings



```
[48]: ptrump_model.wv['cat']
[48]: array([ 0.20400783, -0.039106,  0.01856888,  0.12811816, -0.19469248,
            0.0769702,  0.20283297,  0.29008055,  0.02028482, -0.12926641,
            -0.04846543, -0.05207671, -0.03316664, -0.1765712,  0.01017776,
            0.2708853, -0.08717136,  0.02070806, -0.01205015,  0.14926407,
            0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568,
            -0.0449645, -0.09630641, -0.19584368,  0.06084952,  0.02816491,
            -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,
            -0.03490731,  0.00283731,  0.15054679,  0.4269414,  0.10274042,
            -0.04721209,  0.1454485,  0.24609394, -0.15863296, -0.09804205,
            -0.17660844,  0.1483835, -0.15724996, -0.01741939,  0.06103354,
            -0.00940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,
            -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,
            0.26531997,  0.07995594,  0.04015322,  0.3167902, -0.05698166,
            0.2766356,  0.14080222,  0.07168636, -0.26628318, -0.13518295,
            0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,
            0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,
            0.0483027, -0.11918885,  0.04418642, -0.09205839,  0.03029672,
            0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,
            -0.08515656, -0.19164322,  0.006798,  0.04576558,  0.16265512,
            -0.19076294,  0.37361856, -0.01433535, -0.09413207, -0.12112152],
           dtype=float32)
```

```
[48]: ptrump_model.wv['cat']
[48]: array([ 0.20400783, -0.039106,  0.01856888,  0.12811816, -0.19469248,
            0.0769702,  0.20283297,  0.29008055,  0.02028482, -0.12926641,
            -0.04846543, -0.05207671, -0.03316664, -0.1765712,  0.01017776,
            0.2708853, -0.08717136,  0.02070806, -0.01205015,  0.14926407,
            0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568,
            -0.0449645, -0.09630641, -0.19584368,  0.06084952,  0.02816491,
            -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,
            -0.03490731,  0.00283731,  0.15054679,  0.4269414,  0.10274042,
            -0.04721209,  0.1454485,  0.24609394, -0.15863296, -0.09804205,
            -0.17660844,  0.1483835, -0.15724996, -0.01741939,  0.06103354,
            -0.00940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,
            -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,
            0.26531997,  0.07995594,  0.04015322,  0.3167902, -0.05698166,
            0.2766356,  0.14080222,  0.07168636, -0.26628318, -0.13518295,
            0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,
            0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,
            0.0483027, -0.11918885,  0.04418642, -0.09205839,  0.03029672,
            0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,
            -0.08515656, -0.19164322,  0.006798,  0.04576558,  0.16265512,
            -0.19076294,  0.37361856, -0.01433535, -0.09413207, -0.12112152],
           dtype=float32)
```

```
[48]: ptrump_model.wv['cat']
[48]: array([ 0.20400783, -0.039106,  0.01856888,  0.12811816, -0.19469248,
            0.0769702,  0.20283297,  0.29008055,  0.02028482, -0.12926641,
            -0.04846543, -0.05207671, -0.03316664, -0.1765712,  0.01017776,
            0.2708853, -0.08717136,  0.02070806, -0.01205015,  0.14926407,
            0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568,
            -0.0449645, -0.09630641, -0.19584368,  0.06084952,  0.02816491,
            -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,
            -0.03490731,  0.00283731,  0.15054679,  0.4269414,  0.10274042,
            -0.04721209,  0.1454485,  0.24609394, -0.15863296, -0.09804205,
            -0.17660844,  0.1483835, -0.15724996, -0.01741939,  0.06103354,
            -0.00940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,
            -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,
            0.26531997,  0.07995594,  0.04015322,  0.3167902, -0.05698166,
            0.2766356,  0.14080222,  0.07168636, -0.26628318, -0.13518295,
            0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,
            0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,
            0.0483027, -0.11918885,  0.04418642, -0.09205839,  0.03029672,
            0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,
            -0.08515656, -0.19164322,  0.006798,  0.04576558,  0.16265512,
            -0.19076294,  0.37361856, -0.01433535, -0.09413207, -0.12112152],
           dtype=float32)
```

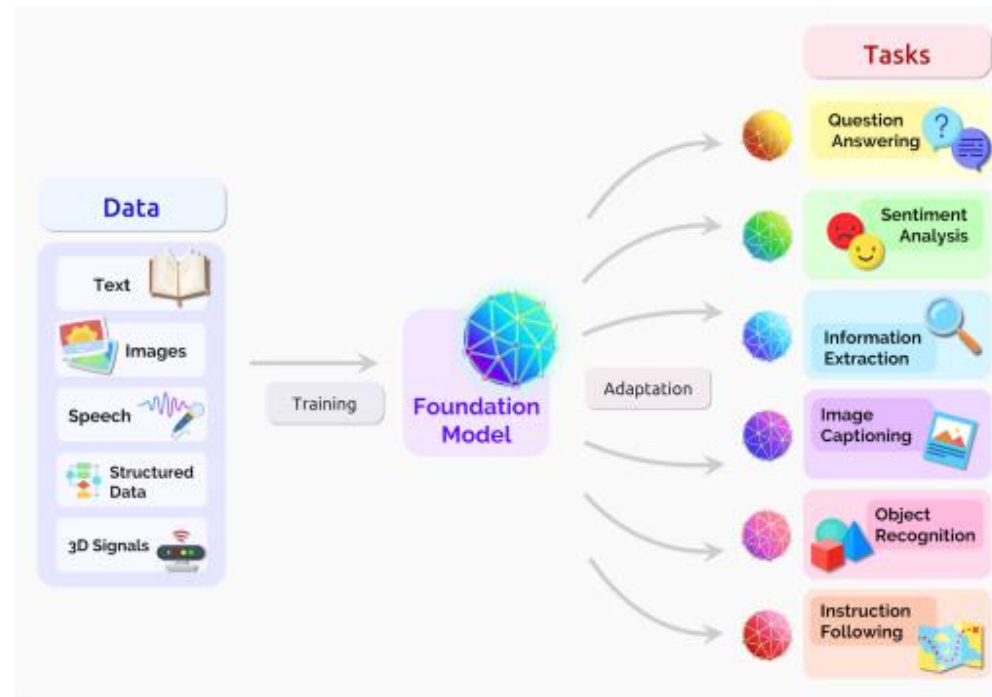
```
[48]: ptrump_model.wv['cat']
[48]: array([ 0.20400783, -0.039106,  0.01856888,  0.12811816, -0.19469248,
            0.0769702,  0.20283297,  0.29008055,  0.02028482, -0.12926641,
            -0.04846543, -0.05207671, -0.03316664, -0.1765712,  0.01017776,
            0.2708853, -0.08717136,  0.02070806, -0.01205015,  0.14926407,
            0.25498936, -0.03342117, -0.13030672,  0.06632984,  0.3417568,
            -0.0449645, -0.09630641, -0.19584368,  0.06084952,  0.02816491,
            -0.12716284, -0.05554081, -0.13319102,  0.09061765, -0.08026118,
            -0.03490731,  0.00283731,  0.15054679,  0.4269414,  0.10274042,
            -0.04721209,  0.1454485,  0.24609394, -0.15863296, -0.09804205,
            -0.17660844,  0.1483835, -0.15724996, -0.01741939,  0.06103354,
            -0.00940341,  0.09271107, -0.08346217,  0.03019234,  0.10981421,
            -0.00283515,  0.03666781, -0.01518192,  0.05991765, -0.01941036,
            0.26531997,  0.07995594,  0.04015322,  0.3167902, -0.05698166,
            0.2766356,  0.14080222,  0.07168636, -0.26628318, -0.13518295,
            0.21819884, -0.00575808, -0.08212868, -0.13839713,  0.08817974,
            0.13988763, -0.10925691, -0.06351396, -0.00679346,  0.01271869,
            0.0483027, -0.11918885,  0.04418642, -0.09205839,  0.03029672,
            0.04010388, -0.12033968,  0.15821207,  0.10007641, -0.07970749,
            -0.08515656, -0.19164322,  0.006798,  0.04576558,  0.16265512,
            -0.19076294,  0.37361856, -0.01433535, -0.09413207, -0.12112152],
           dtype=float32)
```

# Transfer Learning and Foundation Models

# Definitions

**Transfer learning:** Taking knowledge gained from one data set and applying it to another.

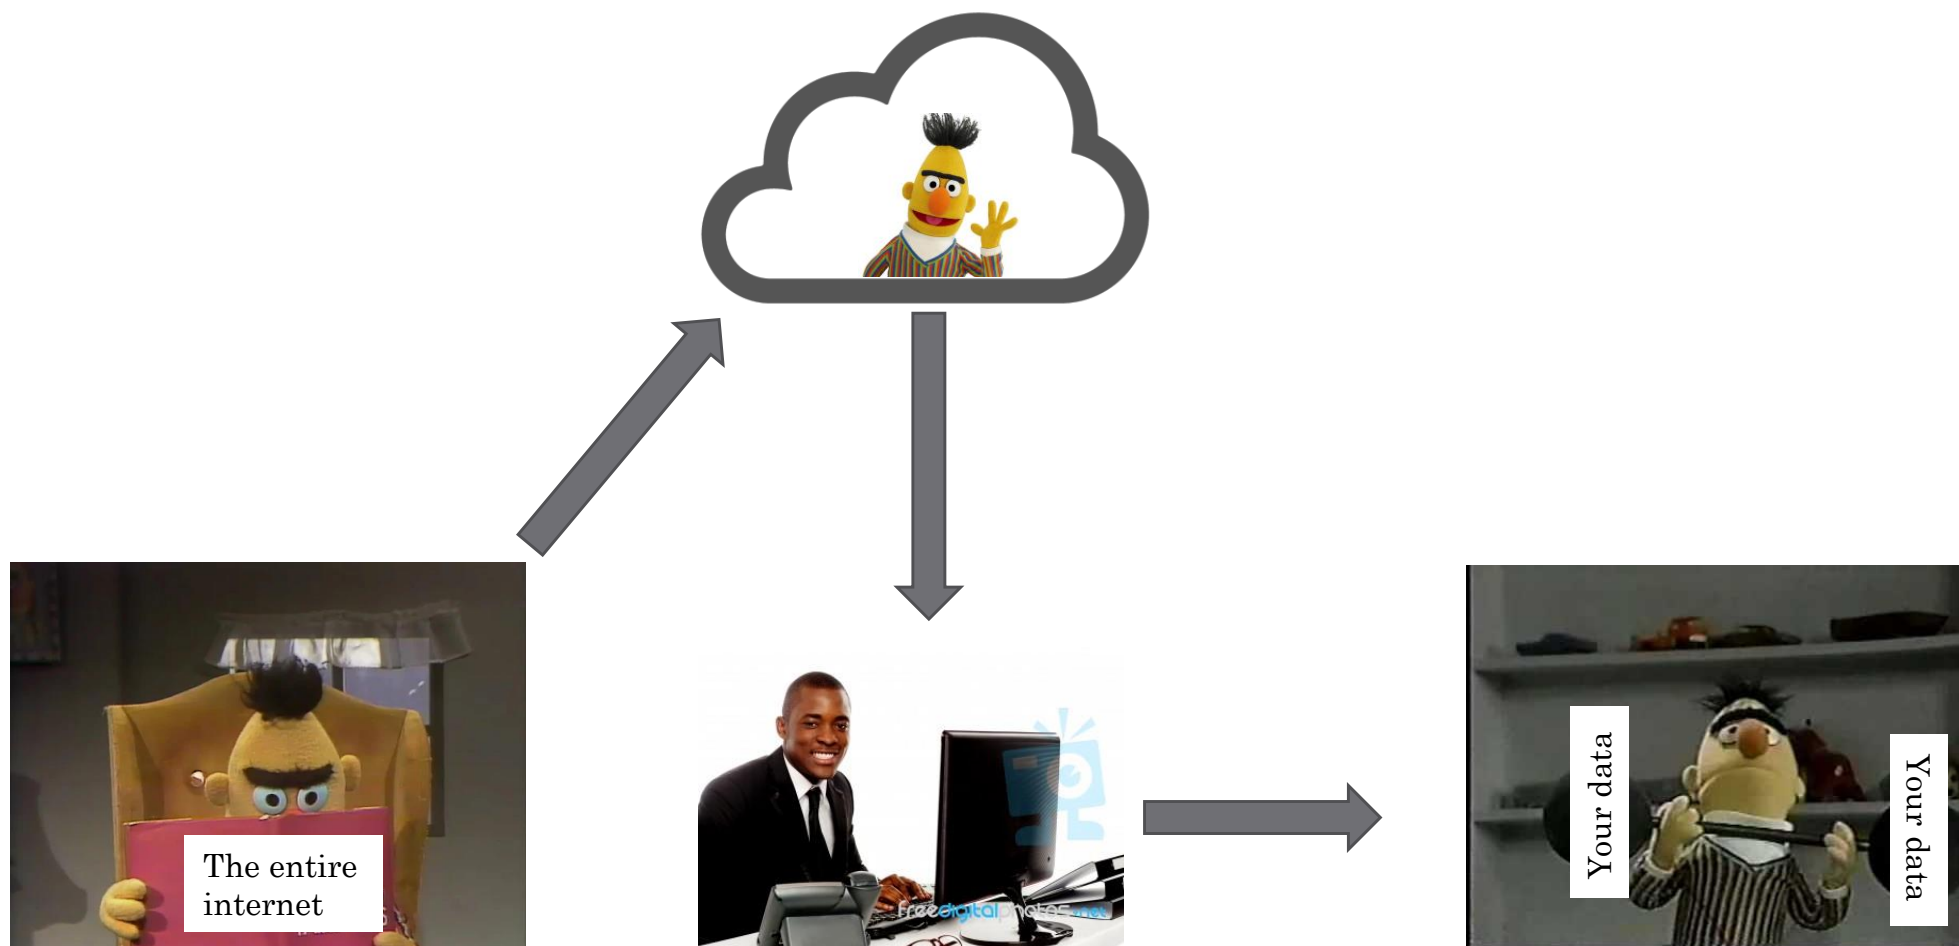
**Foundation Models:** Large models trained on massive data sets intended to be adapted for downstream tasks.



# Bidirectional Encoder Representations from Transformers (BERT)

- Neural network known as a transformer
- Uses contextualized embeddings
- Pre-trained on massive corpora and then adapted to hundreds of thousands of tasks.
- Lots of evolutions:
  - RoBERTa
  - Electra
  - DeBERTa
  - BigBird
  - DistilBERT
  - Etc.

# BERT Supervised Learning Pipeline



# BERT tasks

- Supervised classification
- Zero-shot classification
- Topic models
- Sentiment analysis
- Translation
- Summarization
- Document/sentence similarity
- Question answering

# Generative Pre-trained Transformer (GPT)

- Transformer similar to BERT, but much larger and trained differently
- Primary application is text generation tasks.
- Attractive because of their ease of use
- Currently limited application for social science research because:
  - BERT-like models already match human performance on most tasks researchers care about.
  - Only the largest GPT models can compete with BERT-like models. Both expensive and inefficient.

Classify the following tweets as either support, oppose, or none

###

Tweet: that look when you realize you blew \$350 million for nothing #demdebate #bloomberg

Stance: none

###

Tweet:@USER cool photo from 2012 rnc! idiots! #trump2020

Stance: support

###

Tweet: @USER @USER @USER why you all trying to complicate this???? #votebluetoendthisnightmare #votebluetosaveamerica

Stance: Oppose

###

Tweet: trump is setting up just in case he loses to contest the election. #election2020 #debates #presidentialdebate2020

Stance: oppose



# Hugging Face



# Applications

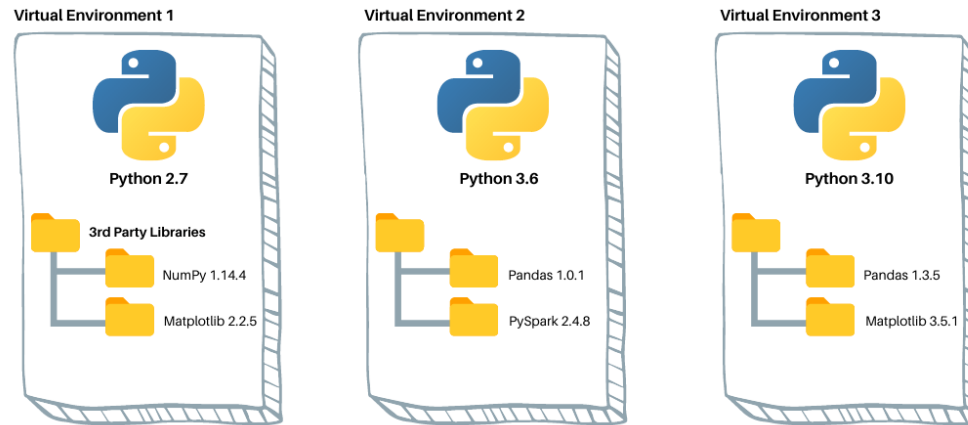
# Some notes on tools: Language



VS.



# Some notes on tools: Environments



[dataquest.io](https://dataquest.io)

# CONDA

# Some notes on tools: Packages

 PyTorch vs.  TensorFlow

 Weights & Biases

# Some notes on tools: Hardware



Or



# Summary:

- Use python if possible
- Always set up a virtual environment. I recommend anaconda.
- Pytorch is generally more user friendly
- Local hardware requires a Nvidia GPU. You can access free/cheap hardware via Google colab

# Coding Demonstration: Classification